



Sharing your data science work:

Transform your R code into an API with plumber

Gustavo Martinez
Francesca Vitalini
Gabriel Foix

An initiative by
Mirai Solutions GmbH
labs@mirai-solutions.com
<https://mirai-solutions.ch>

Bringing ideas to life.

Smart.
Agile.
Personal.



MiraiLabs

MiraiLabs: What is it about?

- Are you a data scientist keen on trying new things?
- Do you use data analysis in your daily work? Do you want to expand your toolkit?
- Are you looking for a more guided hands-on introduction?

Data science workshops for professionals

Data science

- Data analytics
- New tools & techniques
- Interactivity & visualization

For professionals

- Based on real experience from the industry
- Addressing relevant topics

Goals

- Learn together**
- Establish a community

Today's Outline:

- Renku Platform
- Introduction to APIs
- Introduction to plumber
 - Try it yourself section
- Break ~ 10 min
- Plumber by example
 - Try it yourself section
- Best practices and security

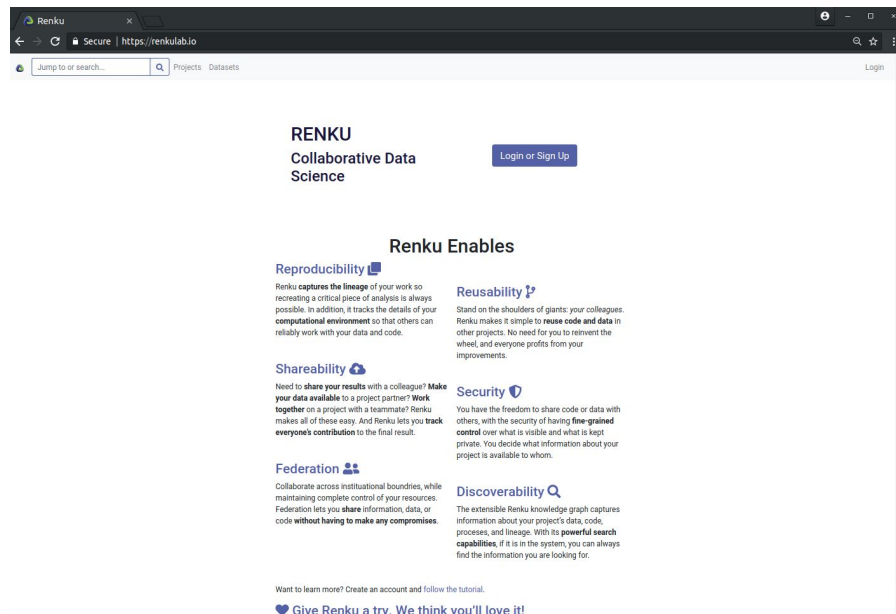
Renku Platform

Renku Platform

Renku is a software platform for data science, developed by the [Swiss Data Science Center](https://www.sdsch.ch/).

The platform is designed to enable reproducible and collaborative data science.

It is an [open source project](https://github.com/renkulab/renku), and also available as a managed service on the cloud, free to use at <https://renkulab.io/>



MiraiLabs on Renku



1.- Open the [MiraiLabs 2.0](https://renkulab.io/projects/gustavo.martinez/mirailabs-2-0/) project and fork it.

The screenshot shows the Renku web interface for the project 'MiraiLabs 2.0' by 'gustavo.martinez/mirailabs-2-0'. The 'fork' button is circled in red. The page includes a search bar, navigation tabs for Overview, Collaboration, Files, Datasets, Environments, and Settings, and a description of the project as a git repository for a workshop.

2.- Log in or register as a new user. You can use your GitHub account to log in.

The screenshot shows the Renku login page. It features a 'Log In' header, a 'Log In' button, and input fields for Email and Password. There are also links for 'Forgot Password?', 'Log In', 'New user?', and 'Register'. The background is a blue and white abstract image.



Spin-up your own environment

The screenshot shows the 'Environments' tab in the MiraiLabs interface. The 'Environments' tab is highlighted with a red box and labeled '1.- Go to the environments tab'. On the left sidebar, the 'New' button is highlighted with a red box and an arrow pointing to it, labeled '2.- New environment'. In the 'Default Environment' section, the '/rstudio' option is highlighted with a red box and labeled '3.- Select rstudio'. At the bottom, the 'Start environment' button is highlighted with a red box and an arrow pointing to it, labeled '4.- Start the environment'.

Overview Collaboration Files Datasets **Environments** Settings

Running

New

Start a new interactive environment

Branch (only 1 available)  

master

Commit  

7598dd89 - Francesca Vitalini - 2020-02-18 15:19:06

Docker Image **available**

Default Environment

/lat **/rstudio**

Number of CPUs

0.25 0.5 1 2

Amount of Memory

1G 2G 4G 8G

Number of GPUs: 0

☐ Automatically fetch LFS data

Start environment




Connect to your environment

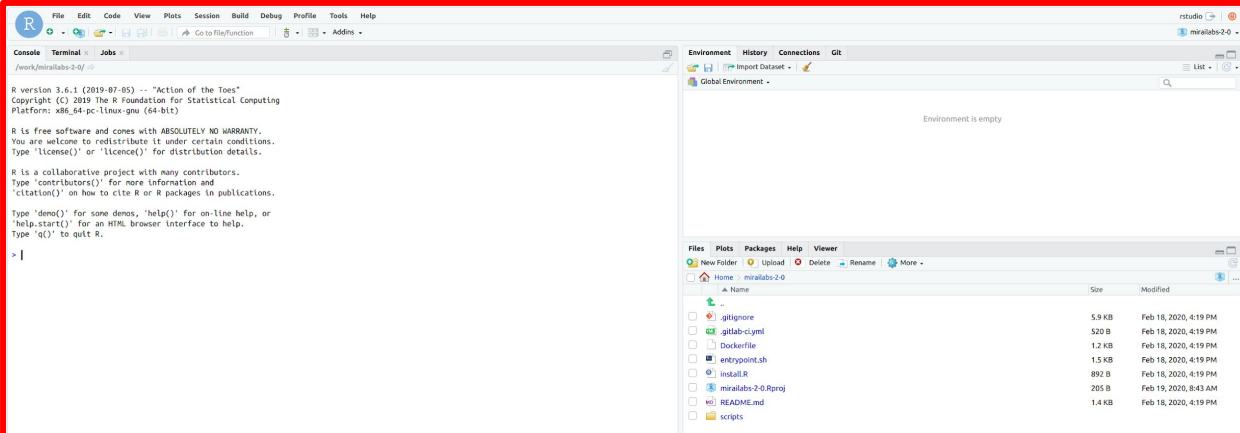
Running

New

Interactive Environments

Branch	Commit	Status	Action
master	7598dd89	Running	Connect





Connect to open up your RStudio session. It will open up the workshop project, with all necessary scripts and dependencies already available/

Web APIs

Web APIs

Web APIs (Application Programming Interface) are communication protocols between computer programs connected in a network.

Unnoticed most of the time, they are the bread and butter of digital life.

Data scientists often use them for data sourcing, often through supporting tools.



HTTP vs REST

[HTTP](#) (Hypertext Transfer Protocol) is a communication protocol for transferring information through a network. It is the most common for web APIs. [HTTPS](#) is a secure form of HTTP.




[REST](#) (REpresentational State Transfer) is a software architectural style for web APIs. REST defines constraints on the software design, not on the technology.



A REST API does not necessarily need to use HTTP.

HTTP




A client sends a request message to a server. The server interprets the request and sends back a response.

HTTP messages are packets of information for which the protocol provides a uniform interface.

HTTP messages



Parts of the request message



Parts of the response message



HTTP message parts

- **URL**¹. The address or endpoint for the request. Can contain parameters.
- **HTTP method**². A specific process invoked on the endpoint.
- **Headers**. Additional information such as who is making the request, what type of response is expected, or format of the message body.
- **Message body**³. Actual data transferred with the message. [Serialized](#).

Commonly used terms referring to the same concepts: ¹URI, ²Verb, ³Payload

HTTP methods

Method	Operation	Description	Status codes
GET	Read	Retrieve a representation of a resource.	200 (OK). Common errors: 404 (NOT FOUND), 400 (BAD REQUEST)
POST	Create	Create new resources, typically subordinated to other (e.g. parent) resource	201 (Created). 409 (Conflict)
PUT	Update / replace	Update a known resource. Sometimes create, if the id of the resource is decided by the client.	200 (OK), 204 (No Content). 404 (Not Found)
DELETE	Delete	Delete a resource	200 (OK). 404 (Not Found), 405 (Method Not Allowed)

Most common methods, used for CRUD operations (Create, Read, Update, Delete)



HTTP status codes

CATEGORY	DESCRIPTION
1xx: Informational	Communicates transfer protocol-level information.
2xx: Success	Indicates that the client's request was accepted successfully.
3xx: Redirection	Indicates that the client must take some additional action in order to complete their request.
4xx: Client Error	This category of error status codes points the finger at clients.
5xx: Server Error	The server takes responsibility for these error status codes.

See more detailed lists in [Wikipedia](#) or [restapitutorial](#)

Web clients

There are multiple options for web clients that can be used to place calls to web APIs and interpret their responses:

- Web browsers addons
- Standalone clients, like [Postman](#) or [Insomnia](#)
- [Swagger UI](#). Interactive UI tool, the choice of RStudio
- Command line tool [curl](#)



Web APIs in R

API wrappers

Many R packages include API wrappers: user-friendly functions for interacting with web APIs from within R.

The wrappers do all the heavy lifting work for us: defining the HTTP requests, and parsing the responses.

For example, the package `tidyquant` integrates resources for collecting financial data from various sources.

Other examples are `bigrquery`, `worldmet` the `cloudyr` project, ...



R web tools

There are multiple packages implementing web clients based on [libcurl](#). We will see some `httr` examples (`01-intro_web_APIs.R`).

`urltools` provides useful tools for handling urls.

`plumber` is a friendly API generator that uses `httpuv`, a low level web server library.

In this workshop we will focus on `plumber`, which is already available in your **Renku** project.

PlumberR


Plumber

`plumber` is an open source R package from RStudio to build and run web APIs. With `plumber`, you can convert existing R code into a web API, by decorating it with special comments.

Get `plumber`:

```
install.packages("plumber")
```

“`plumber`-izing” your R code will make your decorated functions available as API endpoints.

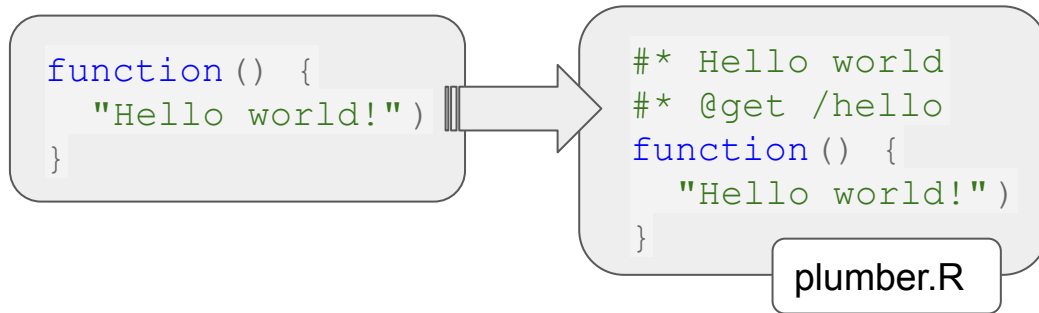


Using plumber: decorating your R code

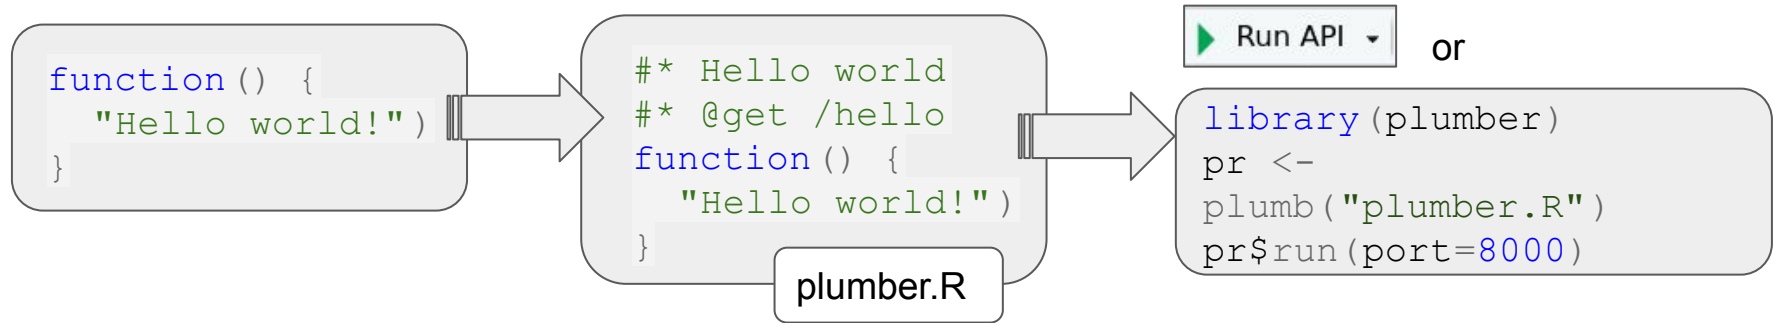


```
function() {  
  "Hello world!"  
}
```

Using plumber: decorating your R code



Using plumber: decorating your R code





Using plumber: Swagger UI

GET /hello Hello world

Parameters
Cancel

No parameters

Execute

Responses

Code	Description	Links
default	Default response.	No links



Responses

Curl

```
curl -X GET "http://127.0.0.1:9817/hello" -H "accept: */*"
```

Request URL

```
http://127.0.0.1:9817/hello
```

Server response

Code	Details
200	<div> Response body <pre>["Hello World"]</pre> Download </div> <div> Response headers <pre>content-length: 15 content-type: application/json date: Thu, 06 Feb 2020 09:41:59 GMT</pre> </div>

Using plumber: serializing the response



By default plumber serializes the response in json format. Other content-types can be specified through annotations:

Annotation	Content Type	Description/References
@json	application/json	jsonlite::toJSON()
@html	text/html; charset=utf-8	Passes response through without any additional serialization
@jpeg	image/jpeg	jpeg()
@png	image/png	png()

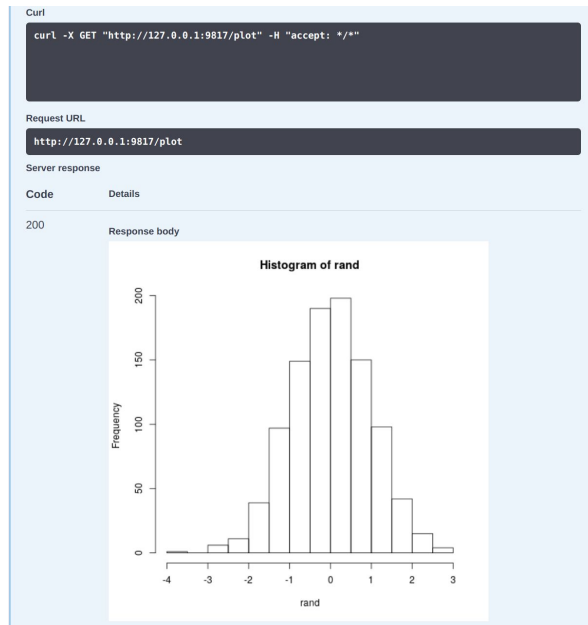


Using plumber: return an image

```
## Plot a histogram
## @png
## @get /plot
function() {
  rand <- rnorm(1000)
  hist(rand)
}
```

The serializer tag `## @png` ensures that the output is correctly interpreted as an image.

Other formats are `## @html` or `## @jpeg`; the default is `## @json`.





Using plumber: passing parameters

```

#* Echo back the input msg
#* @param msg The message to echo
#* @get /echo
function(msg = "") {
  list(msg = paste0("The message
is: '", msg, "'"))
}
    
```

Request URL with parameter in the query string

`http://127.0.0.1:9817/echo?msg=Test`

Curl

```
curl -X GET "http://127.0.0.1:9817/echo?msg=Test" -H "accept: */*"
```

GET /echo Echo back the input msg

Parameters Cancel

Name	Description
msg string (query)	The message to echo

msg - The message to echo

Execute

200

Response body

```
{
  "msg": {
    "The message is: 'Test'"
  }
}
```

Download

Response headers

```
content-length: 34
content-type: application/json
date: Thu, 06 Feb 2020 08:52:17 GMT
```




Using plumber: dynamic routes

```
#* Plot out data from the iris dataset.  
Dynamic route (e.g. 'setosa',  
'virginica')  
#* @param species Filter the data to get  
only this species (e.g. 'setosa')  
#* @get /iris/<species>  
#* @png
```

```
function(species) {  
  ...  
}
```

In a dynamic route, the URL path contains parameters.

The URL combines static parts and parameter values, separated by “/”

Request URL example:

[http://127.0.0.1:9817/iris/ **setosa**](http://127.0.0.1:9817/iris/setosa)



Using plumber: typed dynamic routes

```
/* Return the sum of two numbers.  
Typed dynamic route.  
/* @param a The first number to add  
/* @param b The second number to add  
/* @get /sum/<a:numeric>/<b:numeric>  
function(a,b) {  
  a + b  
}
```

Parameters in typed dynamic routes are the only parameters whose type is controlled by plumber.

Request URL example

`http://127.0.0.1:9817/sum/ 2/1`

R Type	Plumber Name
logical	bool, logical
numeric	double, numeric
integer	int



Using plumber: debugging

Printing messages to track progress:

```
...  
print(paste0("a is of class: ", class(a)))  
...
```

**Interrupt execution for inspection (browser), or debug a function (debug).
Only possible when running interactively:**

```
...  
browser()  
...
```

```
...  
debug(myfunction)  
...
```



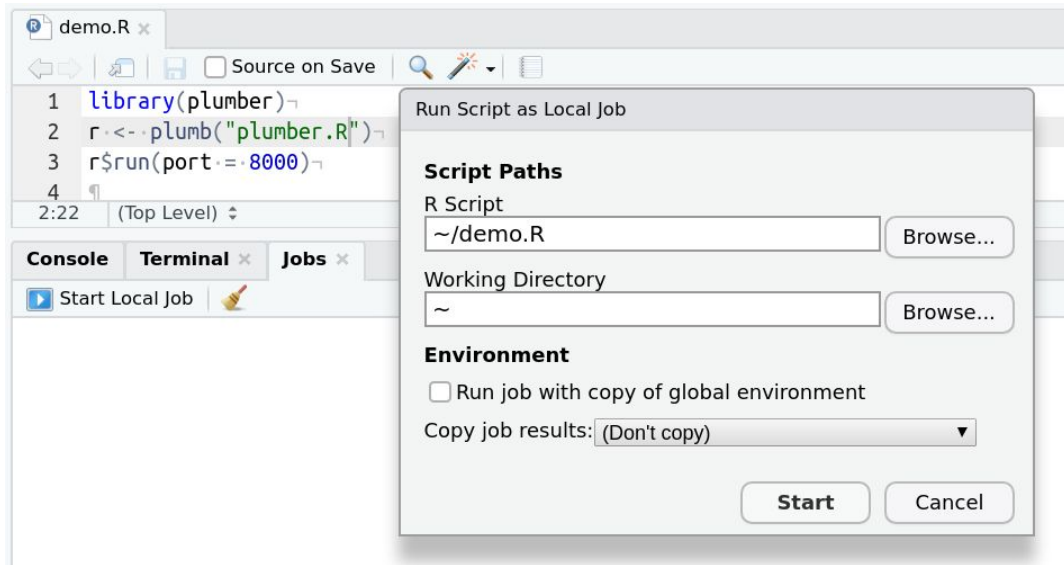
Running plumber as a job in RStudio

Running a `plumber` API as a local job in RStudio, to not have the R session busy.

Multiple jobs can serve multiple APIs at once, in different ports.


```
library(plumber)
pr <-
  plumb("plumber.R")
pr$run(port=8000)
```

demo.R



Plumber resources

- Github: <https://github.com/rstudio/plumber>
- Plumber documentation <https://www.rplumber.io/docs/>
- Webminar <https://rstudio.com/resources/videos/plumbing-apis-with-plumber/>





Try it yourself

- GET current date and time (`Sys.time()`)
- Decorate a function that would return the sentence "Today is ..." based on the input day.
- Consider the function to return the predicted miles per gallon for a given cylinder:

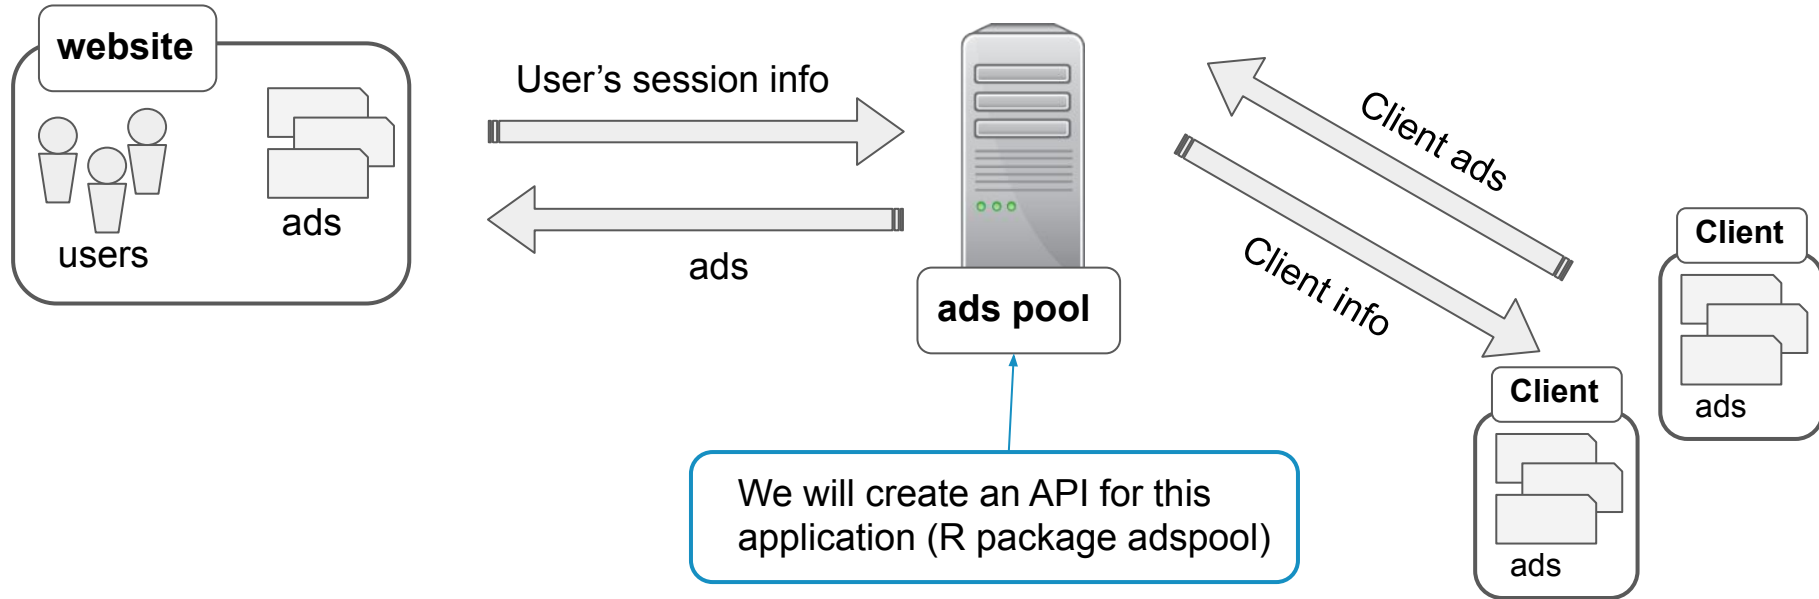
```
function(cyl) {  
  predict(lm(mpg ~ cyl, data = mtcars), data.frame(cyl = cyl))  
}
```

Make a GET endpoint to serve the model (hint: what type should `cyl` have for this to work?)

Plumber

“by example”

The “adverts pool” example



The adspool dataset

The R package, `adspool`, generates the `ads` dataset:

id	name	category	subcategory	img_path	customer_id	click_count	click_rate
1	Mario Kart	Games	Videogame	001_MarioKart.jpeg	10	113	0.02
2	Fifa	Games	Videogame	002_Fifa.jpeg	7	254	0.02
3	Assassin's Creed	Games	Videogame	003_AssassinsCreed.jpeg	6	186	0.03
4	Fortnite	Games	Videogame	004_Fortnite.jpeg	3	120	0.01
5	The Sims	Games	Videogame	005_TheSims.jpeg	9	90	0.02

The adspool tools

The R package, `adspool`, provides tools to:


- Select the best advertisement in a given input context (`select_ads()`)
- Return a subset of advertisements based on input parameters (`subset_ads()`)
- Create an advertisement in the ads pool (`add_ad()`)
- Modify an advertisement in the ads pool (`update_ad()`)
- Remove an advertisement from the ads pool (`remove_ads()`)
- Produce an open invoice for a given customer

The adspool API endpoints

- Get advertisements data for a given name, category, subcategory etc.
- Return all advertisements of a given customer
- Get advertisements profit graph
- Get advertisement image
- Get customer open invoice
- Return ads as R `rds` object
- Delete advertisement from ads pool
- Create a new advertisement in ads pool
- Modify an existing advertisement
- Get advertisements using metadata from a cookie




Plumber routers



- Plumber executes R code in response to HTTP requests.
- Incoming HTTP requests must be routed to one or more R functions.
- Plumber will try to find one (and only one) endpoint with an R function to execute.



Plumber routers



- Filters are functions executed before the endpoint. All filters execute for every HTTP request.
- A filter executes some code and then either returns a response (interrupting execution) or forwards to the next handler (filter or endpoint)
- Both filters and endpoints have access to the request and response objects



The Request object (`req`)

The Request Object in `plumber` is an environment that satisfies the [Rooker Interface](#). `req` encapsulates all the details of the client's request, including:

- Request Path & Method
- Query String, the part of the `URL` after `?`.
- Request Body, called as `req$postBody` (but not exclusive of post methods).
- Cookies, as `req$cookies`.
- HTTP Headers, attached to the request as `req$HTTP_HEADER_NAME`.



The Response object (`res`)

The `plumber` response object `res` is an object of a built-in class, `PlumberResponse`. `res` encapsulates all the information of the response sent back to the client, including:

- Status code
- Body, serialized as per endpoint specification
- Headers

The response object includes some internal methods, notoriously `setCookie`

Plumber cookies



Store information about the state, in a stateless API

```
/* Capture users behavior on a cookie
/* @param usersession
/* @put /ads/usersession/
function(res, usersession) {
  res$setCookie("usersession", usersession)
}
```

`usersession` information available to other endpoints as `req$cookies$usersession`.

Application Examples

- Get advertisements based on browsing history passed as a request body
- Mimic a click on an advertisement
- Get advertisements based on user session (using cookies)
- Log information with a filter logger
- Handle error messages using a filter
- Return a static file

Try it yourself:

Using the `adspool` package as a starting point:

- Write an endpoint to GET all the advertisements belonging to one subcategory. (hint: check `get_customer_ads`)
- Set a `cookie` that will store the time at which the endpoint `get_time` was requested (hint: use `Sys.time()`, check `get_counter()`)
- Create a `filter` that will print to the console the time at which the endpoint was requested (hint: use `Sys.time()`, check `filter_logger()`)

Plumber Advanced

Programmatic usage of plumber

- Creating a router: `pr <- plumber$new()` , object of the class `Plumber`
- Adding endpoints with the method `handle`:

```
pr$handle(methods = "POST",  
           path = "string",  
           handler = function() {},  
           params = list(),  
           serializer = serializer())
```

Add filters and register hooks

- Add filters with the `filter` method:

```
pr$filter(name = "filter name",  
          expr = expiration_in_secs,  
          serializer = serializer())
```

- Define “hooks” to execute code at a particular point of the request’s lifecycle with the `registerHook(s) method(s)`. Currently supported hooks: `preroute`, `postroute`, `preserialize`, `postserialize`

Mounting routers and static files

- Mounting routers with the method `mount`:

```
pr$mount("/path", other_router)
```

Allows modularisation and compartmentalisation of APIs

- Routers can be mounted on top of each other, creating a tree of paths and endpoints
- Static file servers can be combined with routers. These are objects of the class `PlumberStatic`:

```
stat <- PlumberStatic$new("./myfiles")
```

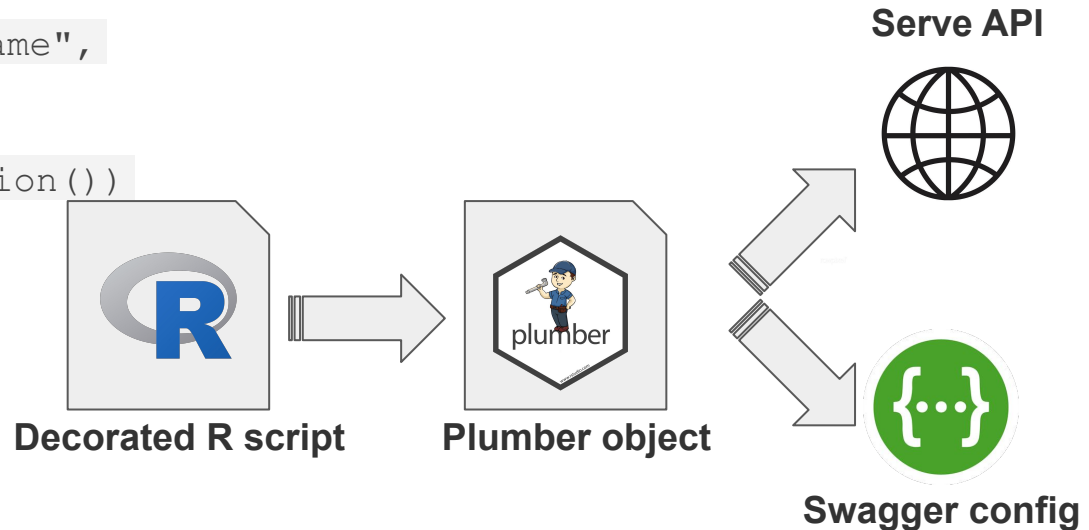
```
pr$mount("/assets", stat)
```

Running plumber

- Serve your API with the `run` method:

```
pr$run(host = "host name",  
      port = 8000,  
      swagger = function())
```

All parameters are optional.



Plumber limitations

- HTTPS not supported natively
- No built-in authentication and no support for `oauth`
- Some common serialization formats are not built-in (e.g. `xml`, `multipart/format`), although one can define custom serializers
- Files upload through endpoint currently [under development](#)
- OpenAPI specs are not always complete. Notoriously: `requestBody` missing
- Functionality exposed through decorators may be limited sometimes

Serving Plumber

- Hosting platforms
 - [DigitalOcean](#)
 - [RStudio Connect](#)
- Hosting on self-managed server
 - [Docker](#)

Serving Plumber. Hosting platforms

- [DigitalOcean](#)

Online, easy to use cloud service, with a small payment plan based on usage. Plumber provides an easy way to deploy an API there (`plumber::do_deploy_api()`)

- [RStudio Connect](#)

Enterprise publishing platform from RStudio. Requires a Licence. It offers buttons to deploy various applications, including APIs. More integration is to be expected as main author of `plumber` also works on RStudio Connect.

Serving Plumber. Self-managed.

A plumber API service is simply a process running an R script. Docker containers provide a convenient way to manage dependencies and ship the final product

- [Docker](#)

Docker image available for plumber: `docker pull trestletech/plumber`

Run your `plumber.R` script (assuming no dependency and that the command is run from the same folder as the script) and serve the API on `localhost:8000/`


```
docker run --rm -p 8000:8000 -v `pwd`/plumber.R:/plumber.R
trestletech/plumber /plumber.R
```

Serving Plumber



**R is single-threaded: one
R session serves one R
process at a time**

Serving Plumber



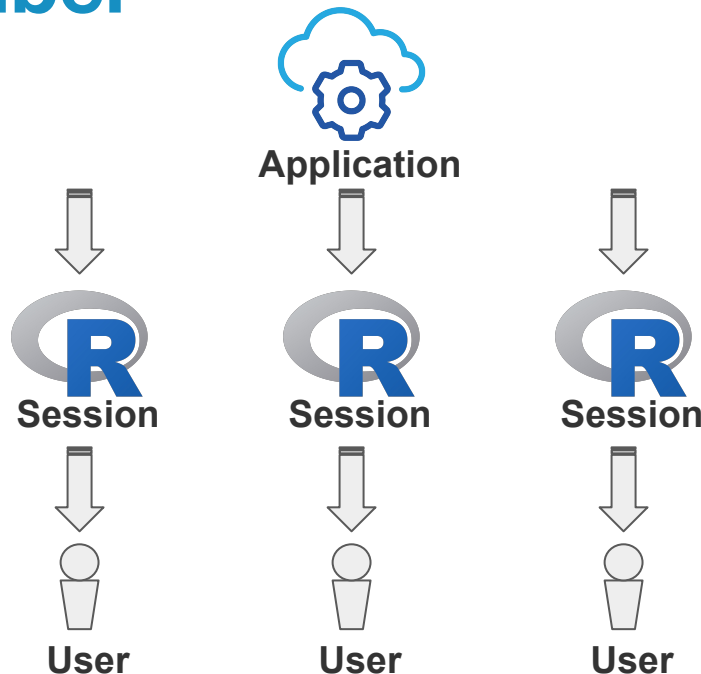
Competing users may queue and have to wait in case of computationally intensive processes

Serving Plumber



**Too many Users.
The application
becomes unusable**

Serving Plumber



A natural solution is serving the API redundantly

This opens new challenges, like load balancing and ensuring data consistency. Good API design practices become critical

Best Practices

Best Practices and Security

Exposing a web API poses challenges that may be new to many R programmers:

- **Best practices.** Commonly accepted standards, and good documentation become critical, in order for the API to be usable by others.
- **Security.** Requires specialized support on a case-by-base. The R programmer needs to be at least aware of some fundamentals.

Best Practices

- Observe generally accepted recommendations for the choice of methods and status codes.
- Control data types.
- Control errors. Return informative error messages, especially for malformed requests.
- Modularize your code. Use router mounting. Make API versions explicit in the URL.
- Use resource expansion and pagination.
- Be as stateless as possible.
- Consider scalability when designing your API.
- Managing state is challenging, leverage transactional data storages when possible.
- Manage I/O: control your connections and close them when no longer necessary.

Security

Beware of the context in which you will expose your API:

- Network environment & firewalls between you and the server hosting the API.
- HTTPS. `plumber` does not support HTTPS natively, this can be added on top though.
- Authentication does not come for free with `plumber` (or any other such framework).

Get familiar with some of the basics of security breaches:

- Code injection
- Denial of Service and Distributed Denial of Service (DoS/DDoS)
- Cross-site scripting



Security. Good practices.

- Manage your secrets: do not store keys/passwords in the code or in places of easy access.
- Restrict the rights of the users that execute your services to the minimum required.
- Distinguish between trusted and untrusted objects. All user input is untrusted: Sanitize.
- Be attentive to the amount of resources your application may consume.
- Think in potential malicious usage of your application.



Q&A

Closing Remarks

MiraiLabs - Next Steps

Feedback welcome:

- Which are the topics you are interested in?
- Format of the workshop: Facilities, materials.

Get in touch!

labs@mirai-solutions.com

Thank you!

Apero Time